

Network Namespace 개요

이름 공간

독립된 공간

가상화기술, Docker등 Container 기술

Virtual Ethernet(veth), pairs

Pipe

Linux Namespace 정의

<http://man7.org/linux/man-pages/man7/namespaces.7.html>

A namespace wraps a global system resource in an abstraction that makes it appear to the processes within **the namespace that they have their own isolated instance of the global resource.** Changes to the global resource are visible to other processes that are members of the namespace, but are invisible to other processes. **One use of namespaces is to implement containers.**

Linux provides the following namespaces:

Namespace	Constant	Isolates
IPC	CLONE_NEWIPC	System V IPC, POSIX message queues
Network	CLONE_NEWNET	Network devices, stacks, ports, etc.
Mount	CLONE_NEWNS	Mount points
PID	CLONE_NEWPID	Process IDs
User	CLONE_NEWUSER	User and group IDs
UTS	CLONE_NEWUTS	Hostname and NIS domain name

Network Namespace 정의

Network namespaces (CLONE_NEWNET)

Network namespaces provide isolation of the system resources associated with networking: network devices, IPv4 and IPv6 protocol stacks, IP routing tables, firewalls, the /proc/net directory, the /sys/class/net directory, port numbers (sockets), and so on.

A physical network device can live in exactly one network namespace.

A virtual network device ("veth") pair provides a pipe-like abstraction that can be used to create tunnels between network namespaces, and can be used to create a bridge to a physical network device in another namespace.

When a network namespace is freed (i.e., when the last process in the namespace terminates), its physical network devices are moved back to the initial network namespace (not to the parent of the process).

Use of network namespaces requires a kernel that is configured with the CONFIG_NET_NS option.

Network Namespace 정의

Global Namespace

Red
Namespace

Net Interfaces,
routing tables,
IPtables

Green
Namespace

Net Interfaces,
routing tables,
IPtables

Net Interfaces, routing tables,
IP stack, IPtables

```
# ip link
# ip address (ip a)
# ip route
# ip netns add red => red namespace 생성
# ip netns add green => green namespace 생성
# ip netns
# ls -l /var/run/netns
```

Global Namespace

Red
Namespace

Net Interfaces,
routing tables,
IPtables

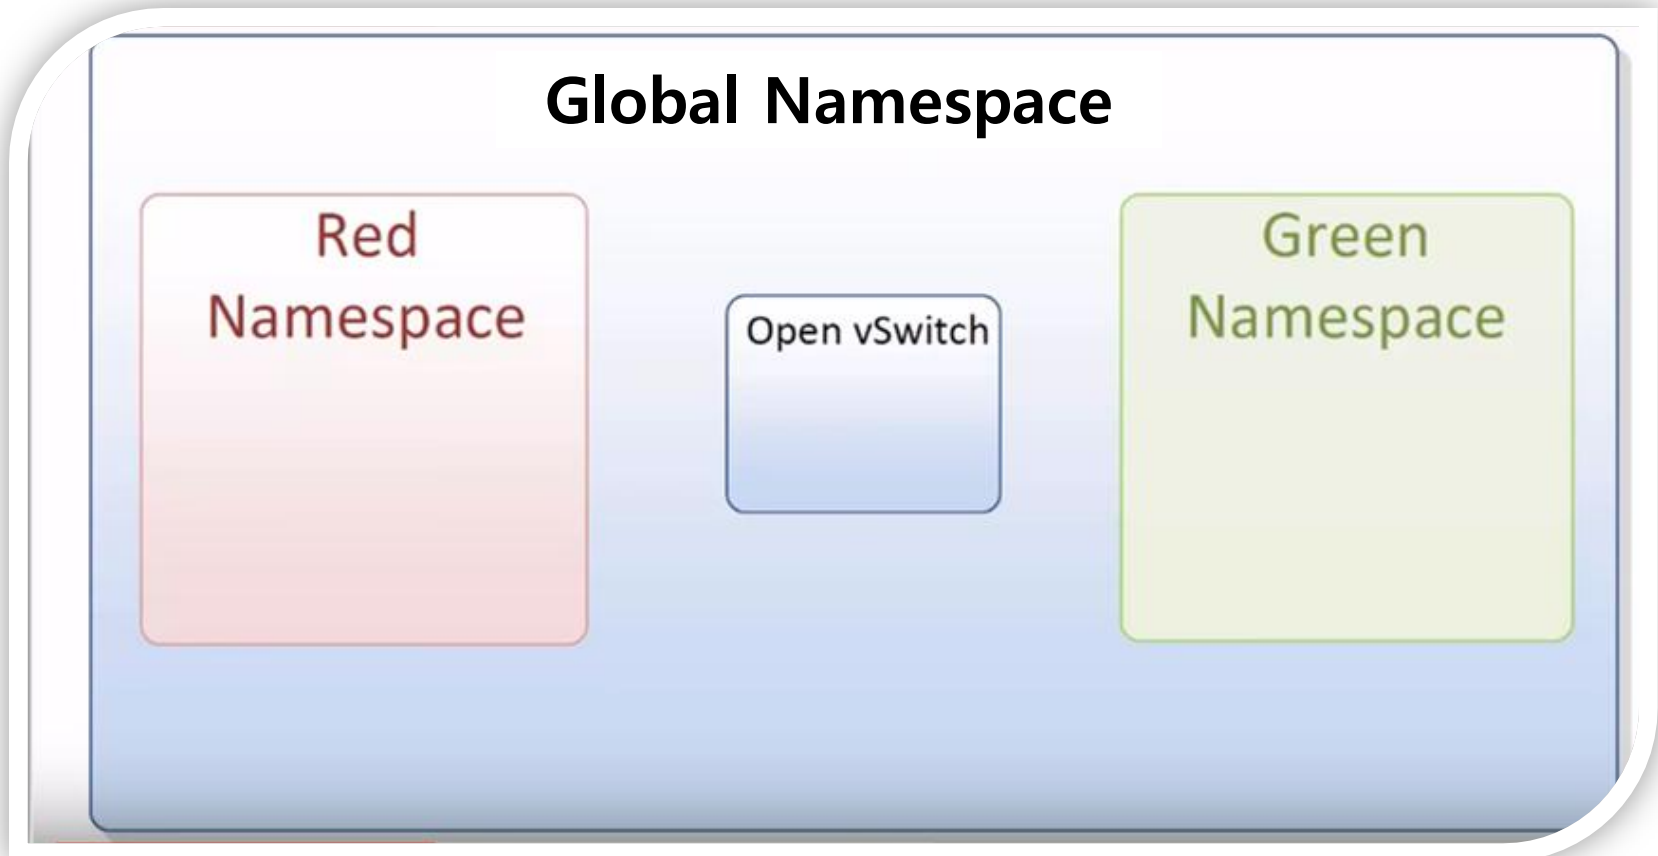
Green
Namespace

Net Interfaces,
routing tables,
IPtables

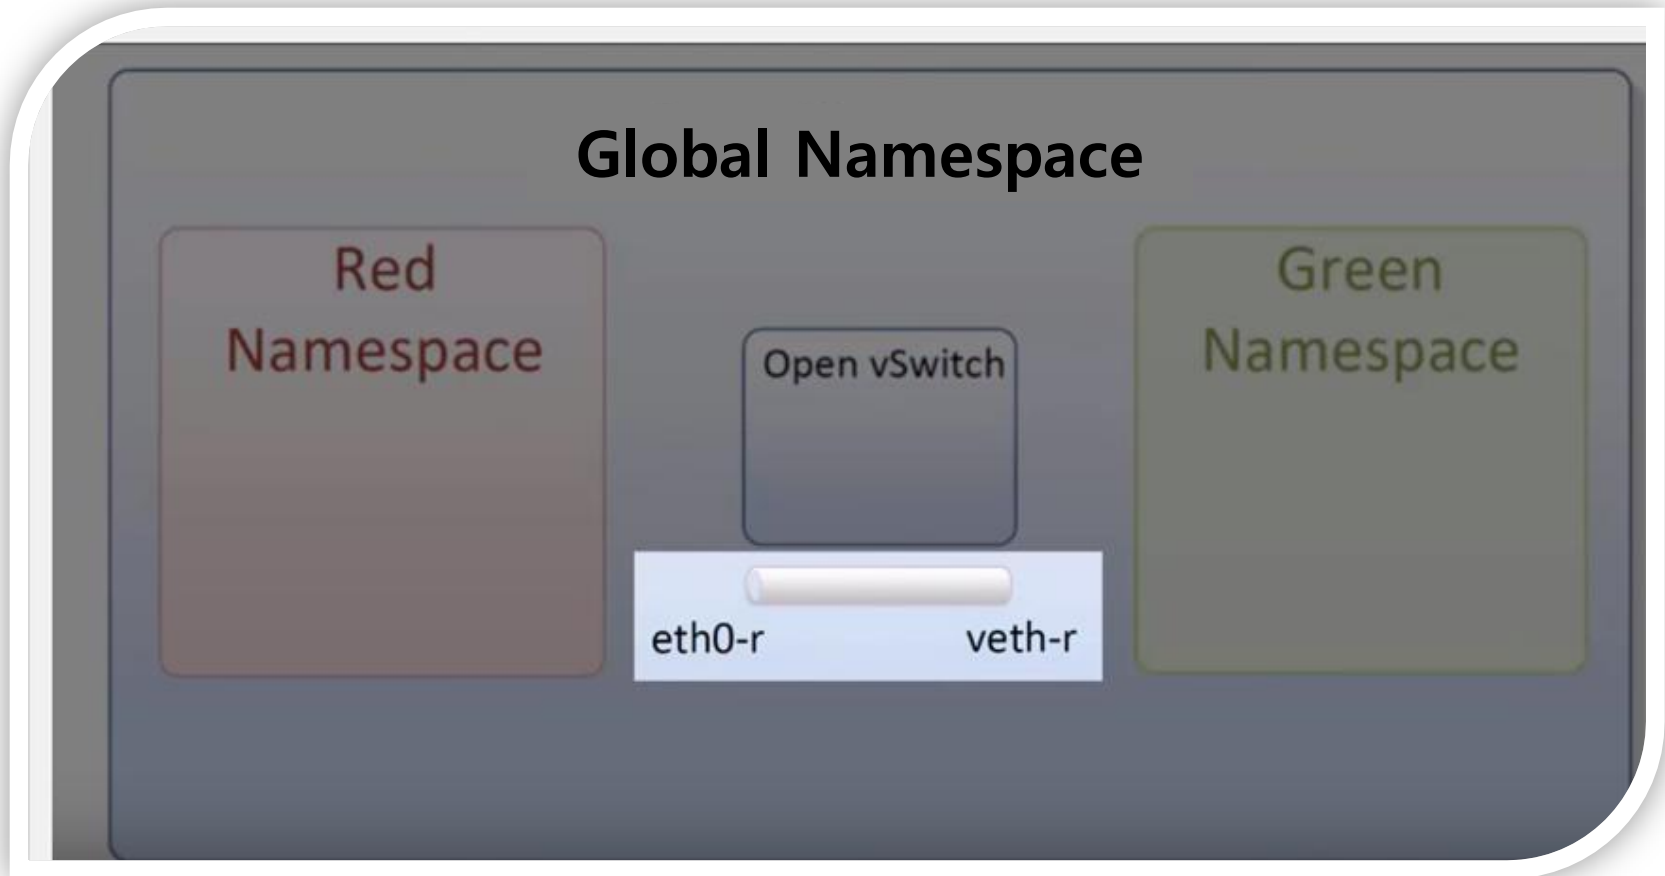
Net Interfaces, routing tables,
IP stack, IPtables

```
# ip netns exec red ip link  
# ip netns exec green ip link
```

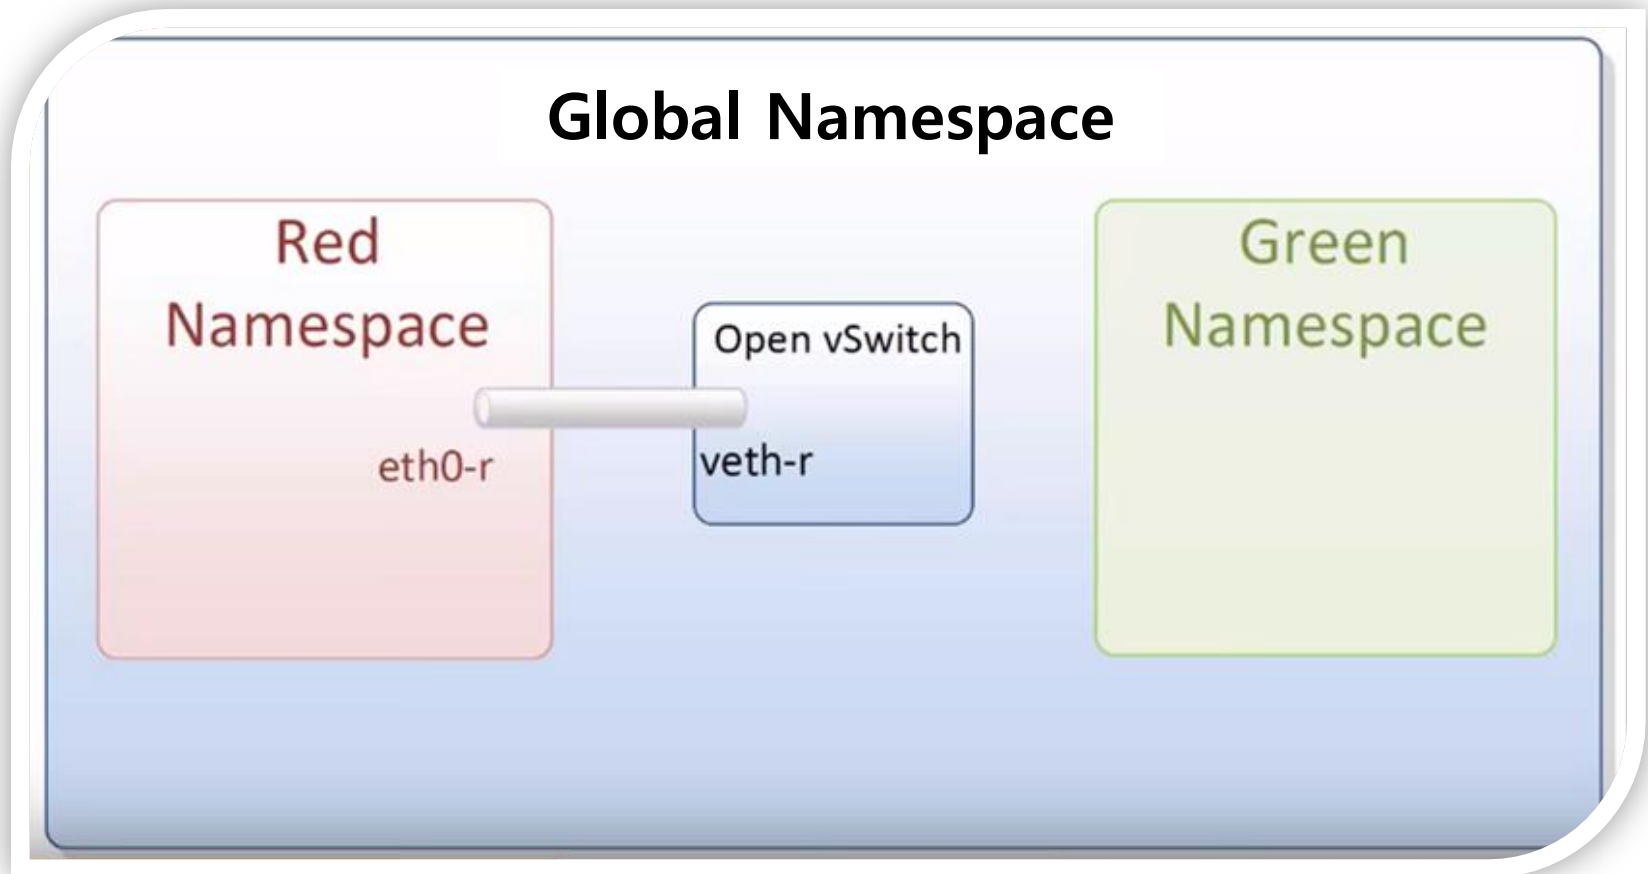
```
# ovs-vsctl add-br OVS1 => Open vSwitch 생성  
# ovs-vsctl show  
# ip link
```



```
# ip link add eth0-r type veth peer name veth-r => veth생성 (Red)
# ip link
# ip link set eth0-r netns red
# ip link
# ip netns exec red ip link
```




ovs-vsctl add-port OVS1 veth-r => OVS내 veth-r 추가
ovs-vsctl show



```
# ip link add eth0-g type veth peer name veth-g => veth생성 (Green)
# ip link
# ip link set eth0-g netns green
# ip netns exec green ip link
# ovs-vsctl add-port OVS1 veth-g => OVS내 veth-g 추가
# ovs-vsctl show
```


Global Namespace



```
# ip link  
# ip link set veth-r up  
# ip link set veth-g up
```

```
# ip netns exec red ip link set dev lo up  
# ip netns exec red ip link set eth0-r up  
# ip netns exec red ip address add 10.0.0.1/24 dev eth0-r  
# ip netns exec red ip a  
# ip netns exec red ip route
```

Global Namespace



ip link set dev veth-g up

ip netns exec green bash

ip link set dev lo up

ip link set eth0-g up


ip address add 10.0.0.2/24 dev eth0-g

ip a

ip route

ping 10.0.0.1; ping 10.0.0.2

Global Namespace



openstack network namespace naming rule

ip netns

qdhcp-[f130c5a0-1879-4e83-91c0-0c9950eb2182, 네트워크 UUID]

: DHCP 프로토콜을 통해 인스턴스에게 IP주소를 할당

snat-[a62685ce-7c40-4836-96f1-0455e332986d, Source NAT UUID]

: 여러 VM이 공용 IP 하나를 공유하여 사용

grouter-[ec565b6f-b0a1-4b41-bfc5-1393b18bc1d9, 라우터 UUID]

: 현재 연결된 서브넷 상의 인스턴스 트래픽 라우팅

qlbaas-[a24f3b84-30d3-4251-8149-092a8696038c, 로드밸런서 UUID]

: 인스턴스에 대한 트래픽을 로드밸런싱하는 Haproxy와 같은 서비스 제공

ip netns help

Usage: ip netns list

ip netns add NAME

ip netns delete NAME

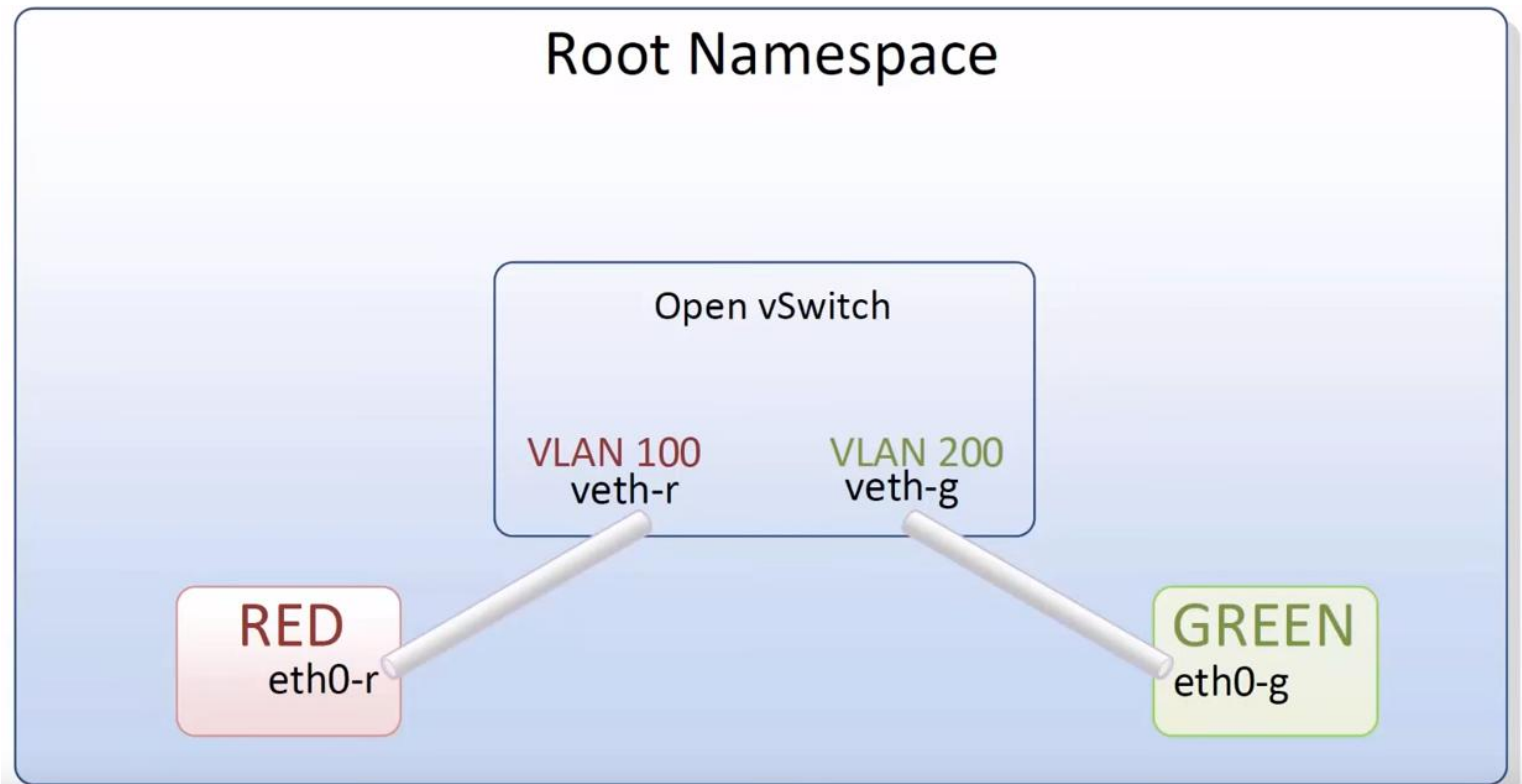
ip netns identify PID

ip netns pids NAME

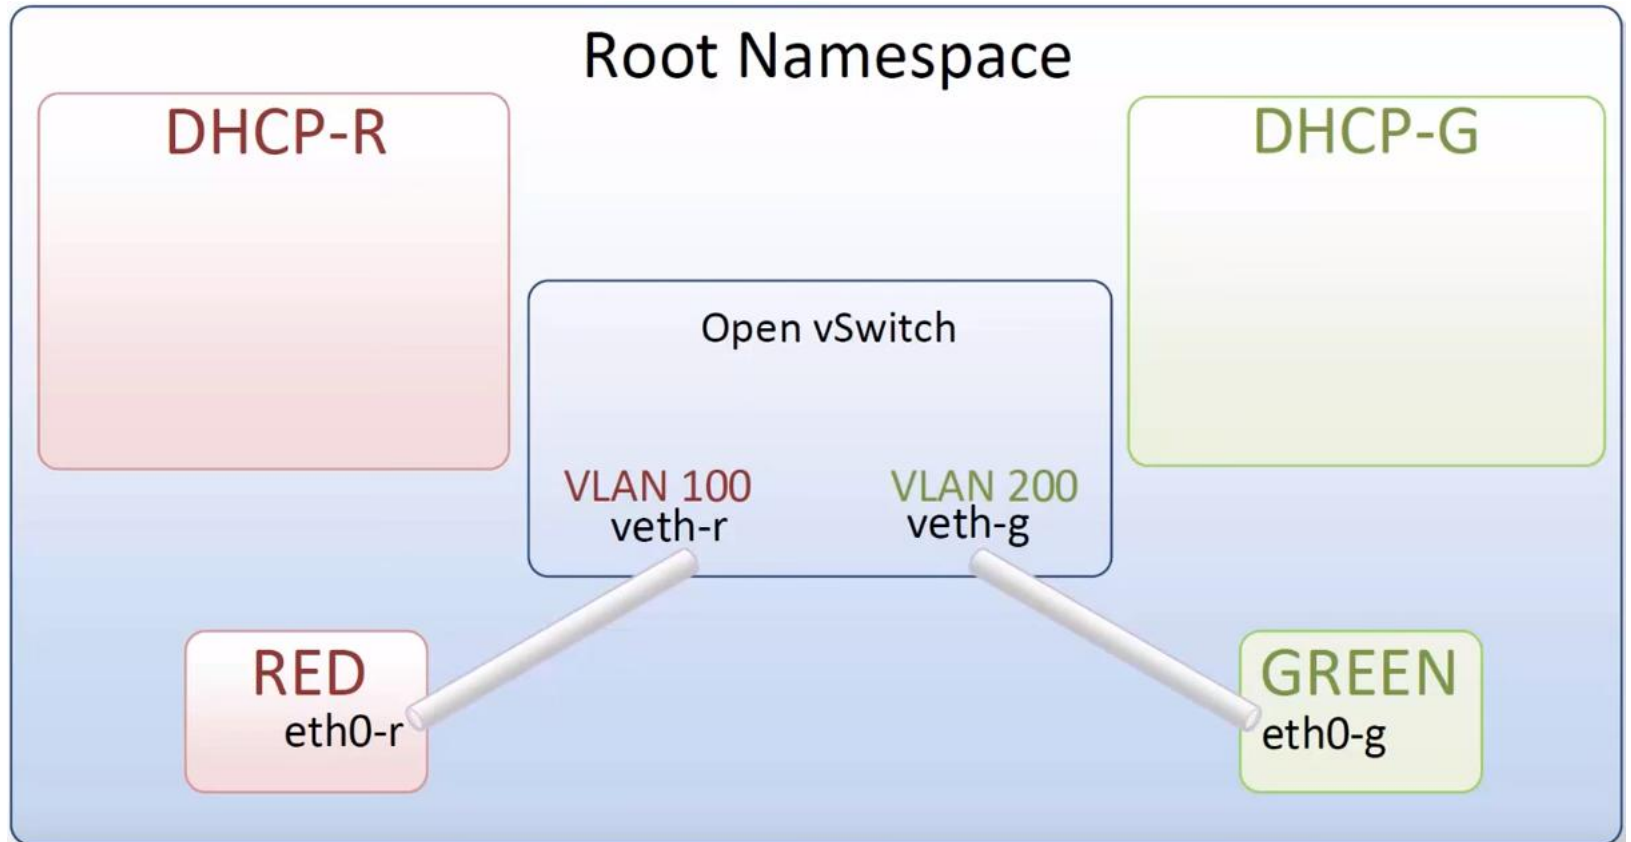
ip netns exec NAME cmd ...

ip netns monitor

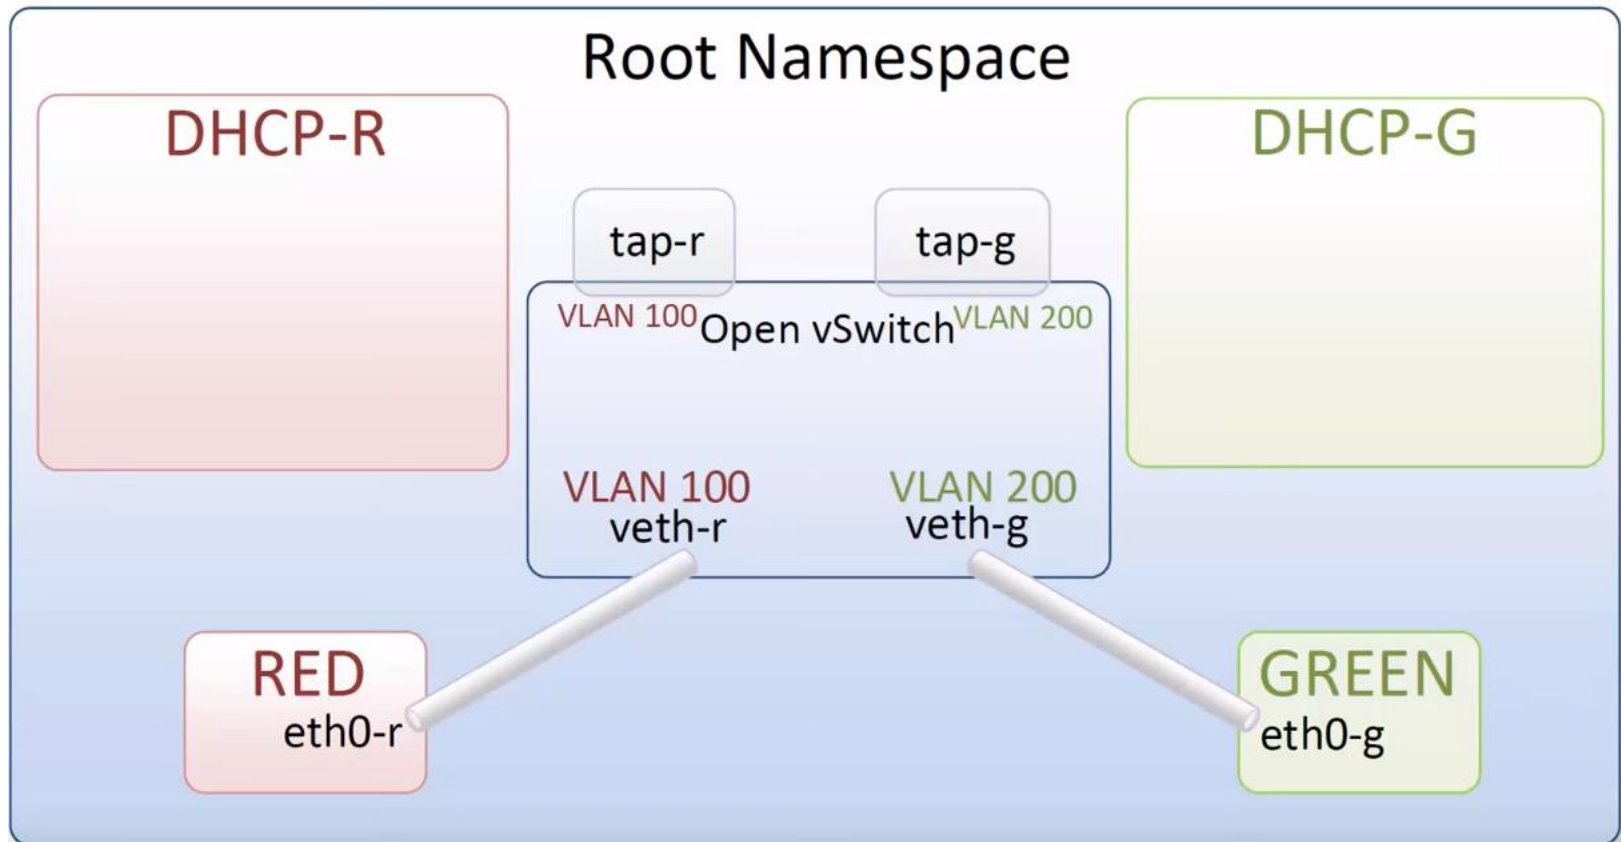
```
# ovs-vsctl set port veth-r tag=100
# ovs-vsctl set port veth-g tag=200
# ovs-vsctl show
# ip netns exec red ip a ; ip netns exec green ip a
# ip netns exec red ip address del 10.0.0.1/24 dev eth0-r
# ip netns exec green ip address del 10.0.0.2/24 dev eth0-g
# ip netns exec red ip a ; ip netns exec green ip a
```




```
# ip netns add dhcp-r ; ip netns add dhcp-g  
# ovs-vsctl show
```



```
# ovs-vsctl add-port OVS1 tap-r
# ovs-vsctl set interface tap-r type=internal
# ovs-vsctl set port tap-r tag=100
# ovs-vsctl add-port OVS1 tap-g
# ovs-vsctl set interface tap-g type=internal
# ovs-vsctl set port tap-g tag=200
```

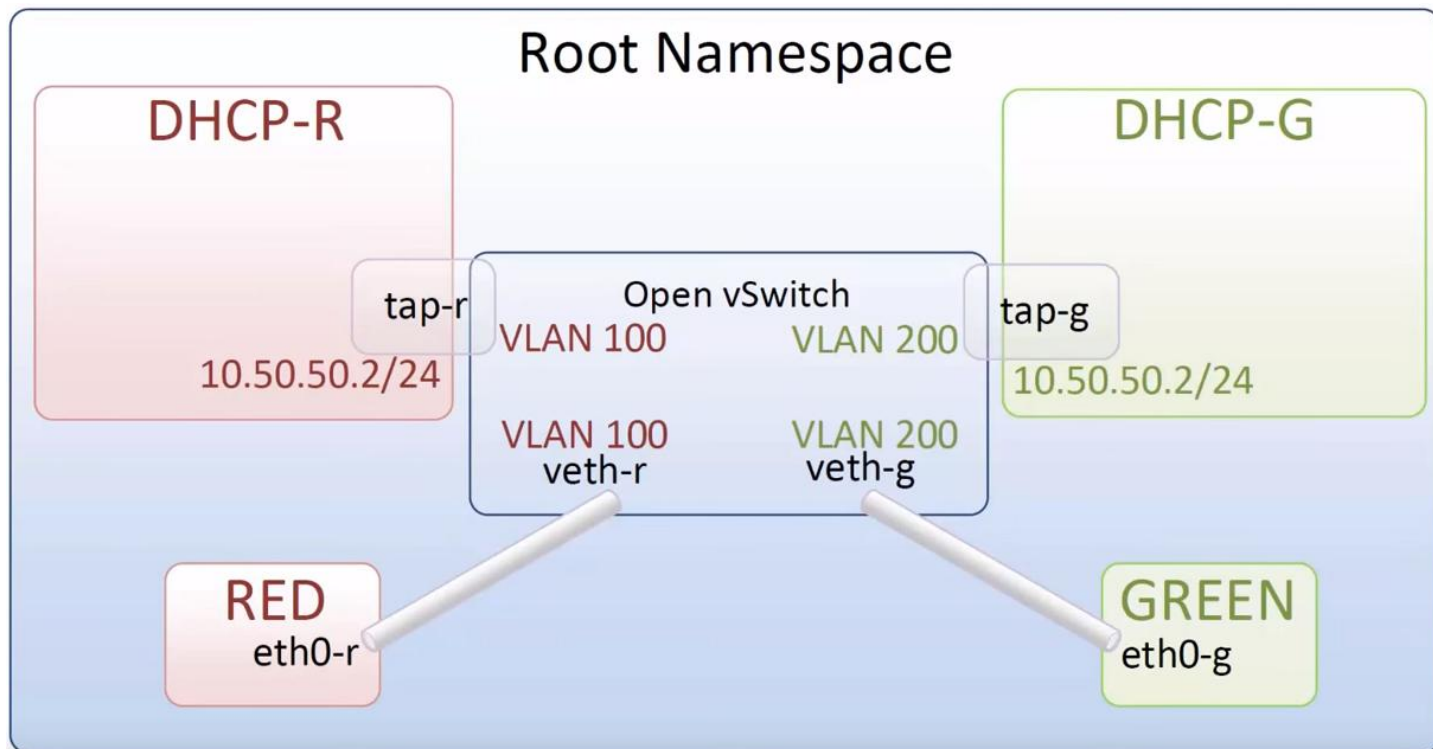


```
# ip netns exec dhcp-r ip link ; ip netns exec dhcp-g ip link
# ip link set tap-r netns dhcp-r
# ip link set tap-g netns dhcp-g
# ip link
# ovs-vsctl show
```



```
# ip netns exec dhcp-r bash
# ip link set dev lo up
# ip link set dev tap-r up
# ip address add 10.50.50.2/24 dev tap-r
```

```
# ip netns exec dhcp-g bash
# ip link set dev lo up
# ip link set dev tap-g up
# ip address add 10.50.50.2/24 dev tap-g
```



ip netns exec dhcp-r dnsmasq --interface=tap-r --dhcp-range=10.50.50.10,10.50.50.100,255.255.255.0

ip netns exec dhcp-g dnsmasq --interface=tap-g --dhcp-range=10.50.50.10,10.50.50.100,255.255.255.0

ps -eaf |grep dns

```
libvirt+ 1590  1 0 10:45 ?      00:00:00 /usr/sbin/dnsmasq --conf-file=/var/lib/libvirt/dnsmasq/default.conf
nobody   2536  833 0 10:46 ?      00:00:00 /usr/sbin/dnsmasq --no-resolv --keep-in-foreground --no-hosts --bind-
interfaces --pid-file=/run/sendsigs.omit.d/network-manager.dnsmasq.pid --listen-address=127.0.1.1 --conf-
file=/var/run/NetworkManager/dnsmasq.conf --cache-size=0 --proxy-dnssec --enable-
dbus=org.freedesktop.NetworkManager.dnsmasq --conf-dir=/etc/NetworkManager/dnsmasq.d
nobody   4295  1 0 18:01 ?      00:00:00 dnsmasq --interface=tap-r --dhcp-
range=10.50.50.10,10.50.50.100,255.255.255.0
nobody   4368  1 0 18:01 ?      00:00:00 dnsmasq --interface=tap-g --dhcp-
range=10.50.50.10,10.50.50.100,255.255.255.0
root     4730  7732 0 18:04 pts/25  00:00:00 grep --color=auto dns
root@stack-VirtualBox:~#
```

ip netns identify 4295

dhcp-r

ip netns identify 4368

dhcp-g

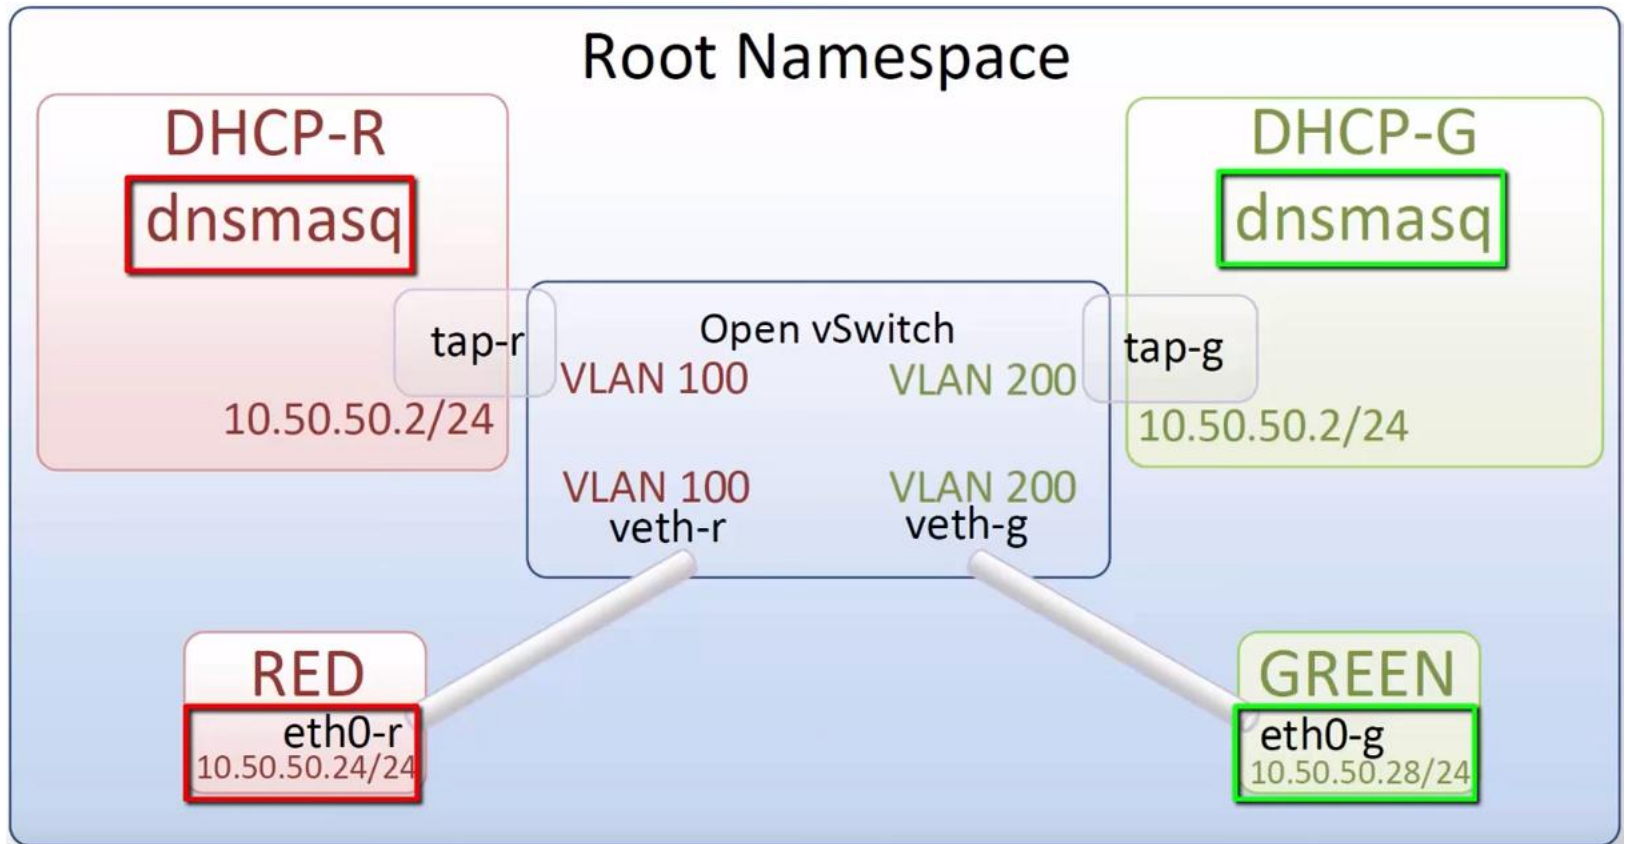
ip netns exec red dhclient eth0-r

ip netns exec red ip a

ip netns exec green dhclient eth0-g

ip netns exec green ip a

#



참고

1. https://www.youtube.com/watch?v=_WgUwUf1d34
2. <http://man7.org/linux/man-pages/man8/ip-netns.8.html>
3. <http://man7.org/linux/man-pages/man7/namespaces.7.html>
4. <http://man7.org/linux/man-pages/man2/clone.2.html>
5. <http://man7.org/linux/man-pages/man2/setns.2.html>
6. <http://man7.org/linux/man-pages/man2/unshare.2.html>